

Importing the Libraries ✓

```
In [1]: #We start off this project by importing all the necessary libraries that will be required for the process
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Loading the Data ✓

```
In [2]: #Loading the data and removing the irrelevant columns.
df = pd.read_csv('smoking.csv')
df.head()
```

Out[2]:

weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	...	hemoglobin	Urine protein	serum creatinine	AST	ALT	Gtp	oral	dental caries	tartar	smoking
60	81.3	1.2	1.0	1.0	1.0	...	12.9	1.0	0.7	18.0	19.0	27.0	Y	0	Y	0
60	81.0	0.8	0.6	1.0	1.0	...	12.7	1.0	0.6	22.0	19.0	18.0	Y	0	Y	0
60	80.0	0.8	0.8	1.0	1.0	...	15.8	1.0	1.0	21.0	16.0	22.0	Y	0	N	1
70	88.0	1.5	1.5	1.0	1.0	...	14.7	1.0	1.0	19.0	26.0	18.0	Y	0	Y	0
60	86.0	1.0	1.0	1.0	1.0	...	12.5	1.0	0.6	16.0	14.0	22.0	Y	0	N	0

```
In [3]: df = df.drop(columns=['ID', 'oral']) ✓  
df.head()
```

Out[3]:

waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	...	LDL	hemoglobin	Urine protein	serum creatinine	AST	ALT	Gtp	dental caries	tartar	smoking
81.3	1.2	1.0	1.0	1.0	114.0	...	126.0	12.9	1.0	0.7	18.0	19.0	27.0	0	Y	0
81.0	0.8	0.6	1.0	1.0	119.0	...	127.0	12.7	1.0	0.6	22.0	19.0	18.0	0	Y	0
80.0	0.8	0.8	1.0	1.0	138.0	...	151.0	15.8	1.0	1.0	21.0	16.0	22.0	0	N	1
88.0	1.5	1.5	1.0	1.0	100.0	...	226.0	14.7	1.0	1.0	19.0	26.0	18.0	0	Y	0
86.0	1.0	1.0	1.0	1.0	120.0	...	107.0	12.5	1.0	0.6	16.0	14.0	22.0	0	N	0

```
In [4]: #Checking the shape of a dataframe and datatypes of all columns along with calculating the statistical data. ✓  
df.shape
```

Out[4]: (55692, 25)

In [5]: df.info() ✓

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55692 entries, 0 to 55691
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   gender                                55692 non-null  object
1   age                                    55692 non-null  int64
2   height(cm)                           55692 non-null  int64
3   weight(kg)                            55692 non-null  int64
4   waist(cm)                             55692 non-null  float64
5   eyesight(left)                        55692 non-null  float64
6   eyesight(right)                       55692 non-null  float64
7   hearing(left)                         55692 non-null  float64
8   hearing(right)                        55692 non-null  float64
9   systolic                             55692 non-null  float64
10  relaxation                            55692 non-null  float64
11  fasting blood sugar                   55692 non-null  float64
12  Cholesterol                           55692 non-null  float64
13  triglyceride                          55692 non-null  float64
14  HDL                                    55692 non-null  float64
15  LDL                                    55692 non-null  float64
16  hemoglobin                            55692 non-null  float64
17  Urine protein                         55692 non-null  float64
18  serum creatinine                      55692 non-null  float64
19  AST                                    55692 non-null  float64
20  ALT                                    55692 non-null  float64
21  Gtp                                    55692 non-null  float64
22  dental caries                          55692 non-null  int64
23  tartar                                55692 non-null  object
24  smoking                               55692 non-null  int64
dtypes: float64(18), int64(5), object(2)
memory usage: 10.6+ MB
```

In [6]: df.describe() ✓

Out[6]:

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	relaxation	...	
count	55692.000000	55692.000000	55692.000000	55692.000000	55692.000000	55692.000000	55692.000000	55692.000000	55692.000000	55692.000000	...	55692.000000
mean	44.182917	164.649321	65.864936	82.046418	1.012623	1.007443	1.025587	1.026144	121.494218	76.004830	...	76.004830
std	12.071418	9.194597	12.820306	9.274223	0.486873	0.485964	0.157902	0.159564	13.675989	9.679278	...	9.679278
min	20.000000	130.000000	30.000000	51.000000	0.100000	0.100000	1.000000	1.000000	71.000000	40.000000	...	40.000000
25%	40.000000	160.000000	55.000000	76.000000	0.800000	0.800000	1.000000	1.000000	112.000000	70.000000	...	70.000000
50%	40.000000	165.000000	65.000000	82.000000	1.000000	1.000000	1.000000	1.000000	120.000000	76.000000	...	76.000000
75%	55.000000	170.000000	75.000000	88.000000	1.200000	1.200000	1.000000	1.000000	130.000000	82.000000	...	82.000000
max	85.000000	190.000000	135.000000	129.000000	9.900000	9.900000	2.000000	2.000000	240.000000	146.000000	...	146.000000

8 rows × 23 columns




Missing Values ✓

```
In [7]: #Checking out the missing values in a dataframe  
df.isnull().sum()
```


```
Out[7]: gender                0  
age                0  
height(cm)         0  
weight(kg)          0  
waist(cm)           0  
eyesight(left)      0  
eyesight(right)     0  
hearing(left)       0  
hearing(right)      0  
systolic            0  
relaxation          0  
fasting blood sugar 0  
Cholesterol         0  
triglyceride        0  
HDL                 0  
LDL                 0  
hemoglobin          0  
Urine protein       0  
serum creatinine    0  
AST                 0  
ALT                 0  
Gtp                 0  
dental caries       0  
tartar              0  
smoking             0  
dtype: int64
```

Data Visualization ✓

```
In [12]: #We can clearly see from the below graph that most smokers are men  
sns.barplot(x=df['gender'],y=df['smoking'])  
plt.show()
```




```
In [14]: sns.countplot(x='gender', hue='smoking', data=df)
plt.show()
```



```
In [15]: #There are 36.73 percent of the people who are smoking a cigarette.
```


```
In [15]: #There are 36.73 percent of the people who are smoking cigarette. ✓  
plt.figure(figsize=(10,5))  
df['smoking'].value_counts().plot.pie(autopct='%0.2f')
```

```
Out[15]: <Axes: ylabel='smoking'>
```





```
In [16]: #Most number of smokers are having the age 40 ✓  
plt.figure(figsize=(9,6))  
sns.histplot(x='age', hue='smoking', data=df)
```


Out[16]: <Axes: xlabel='age', ylabel='Count'>



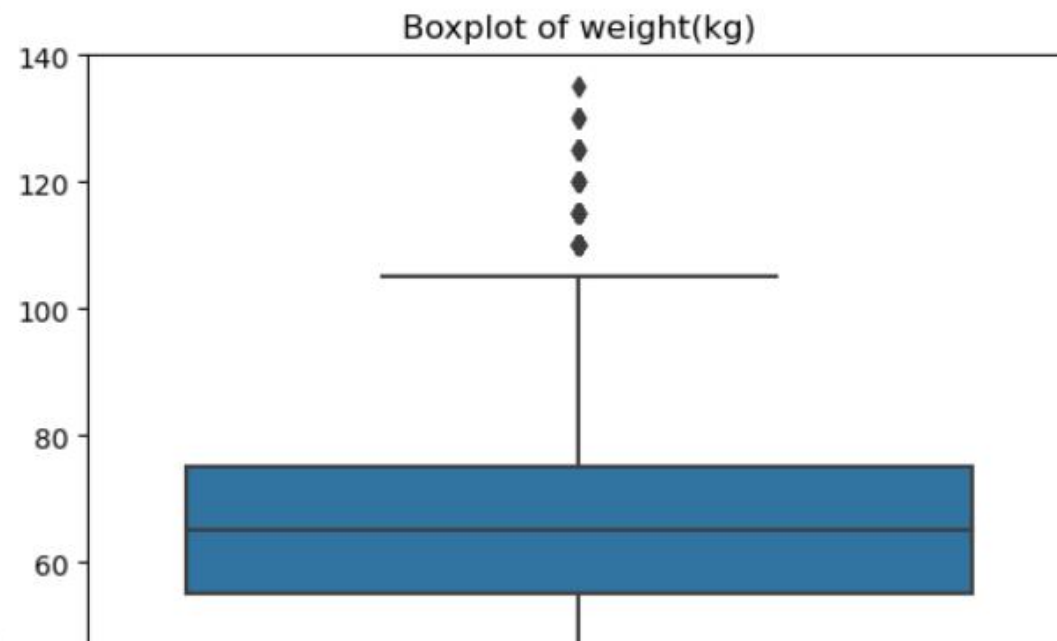
```
In [21]: #Representation of columns using boxplot to detect outliers. Here outliers represent natural
#variations in the population, and they should be left as is in the dataset. These are called true
#outliers. Therefore for this dataset we will not remove outlier.
for i in df.columns:
    if(df[i].dtypes=="int64" or df[i].dtypes=='float64'):
        sns.boxplot(df[i]).set(title=f'Boxplot of {i}')
        plt.show()
```




```
In [21]: #Representation of columns using boxplot to detect outliers. Here outliers represent natural  
#variations in the population, and they should be left as is in the dataset. These are called true  
#outliers. Therefore for this dataset we will not remove outlier.  
for i in df.columns:  
    if(df[i].dtypes=="int64" or df[i].dtypes=='float64'):  
        sns.boxplot(df[i]).set(title=f'Boxplot of {i}')  
        plt.show()
```



```
In [21]: #Representation of columns using boxplot to detect outliers. Here outliers represent natural  
#variations in the population, and they should be left as is in the dataset. These are called true  
#outliers. Therefore for this dataset we will not remove outlier.  
for i in df.columns:  
    if(df[i].dtypes=="int64" or df[i].dtypes=='float64'):  
        sns.boxplot(df[i]).set(title=f'Boxplot of {i}')  
        plt.show()
```




```
In [21]: #Representation of columns using boxplot to detect outliers. Here outliers represent natural
#variations in the population, and they should be left as is in the dataset. These are called true
#outliers. Therefore for this dataset we will not remove outlier.
for i in df.columns:
    if(df[i].dtypes=="int64" or df[i].dtypes=="float64"):
        sns.boxplot(df[i]).set(title=f'Boxplot of {i}')
        plt.show()
```



In [21]: *#Representation of columns using boxplot to detect outliers. Here outliers represent natural variations in the population, and they should be left as is in the dataset. These are called true outliers. Therefore for this dataset we will not remove outlier.*

```
for i in df.columns:  
    if(df[i].dtypes=="int64" or df[i].dtypes=="float64"):  
        sns.boxplot(df[i]).set(title=f'Boxplot of {i}')  
        plt.show()
```



Data Cleaning

In [23]: *#Performing One Hot Encoding for categorical features of a dataframe*

```
from sklearn.preprocessing import LabelEncoder  
|  
le= LabelEncoder()  
df["gender"]=le.fit_transform(df["gender"])  
df["tartar"]=le.fit_transform(df["tartar"])  
  
df["dental caries"]=le.fit_transform(df["dental caries"])  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 55692 entries, 0 to 55691  
Data columns (total 25 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   gender                                55692 non-null  int64  
1   age                                   55692 non-null  int64  
2   height(cm)                           55692 non-null  int64  
3   weight(kg)                            55692 non-null  int64  
4   waist(cm)                             55692 non-null  float64  
5   eyesight(left)                        55692 non-null  float64  
6   eyesight(right)                       55692 non-null  float64  
7   hearing(left)                         55692 non-null  float64  
8   hearing(right)                        55692 non-null  float64  
9   systolic                              55692 non-null  float64  
10  relaxation                             55692 non-null  float64  
11  fasting blood sugar                    55692 non-null  float64  
12  Cholesterol                            55692 non-null  float64  
13  triglyceride                           55692 non-null  float64  
14  HDL                                    55692 non-null  float64  
15  LDL                                    55692 non-null  float64  
16  hemoglobin                             55692 non-null  float64  
17  Urine protein                          55692 non-null  float64  
18  serum creatinine                       55692 non-null  float64  
19  AST                                    55692 non-null  float64  
20  ALT                                    55692 non-null  float64  
21  Gtp                                     55692 non-null  float64  
22  dental caries                           55692 non-null  int64  
23  tartar                                 55692 non-null  int64  
24  smoking                                55692 non-null  int64  
dtypes: float64(18), int64(7)  
memory usage: 10.6 MB
```

Feature selection using feature importance ✓

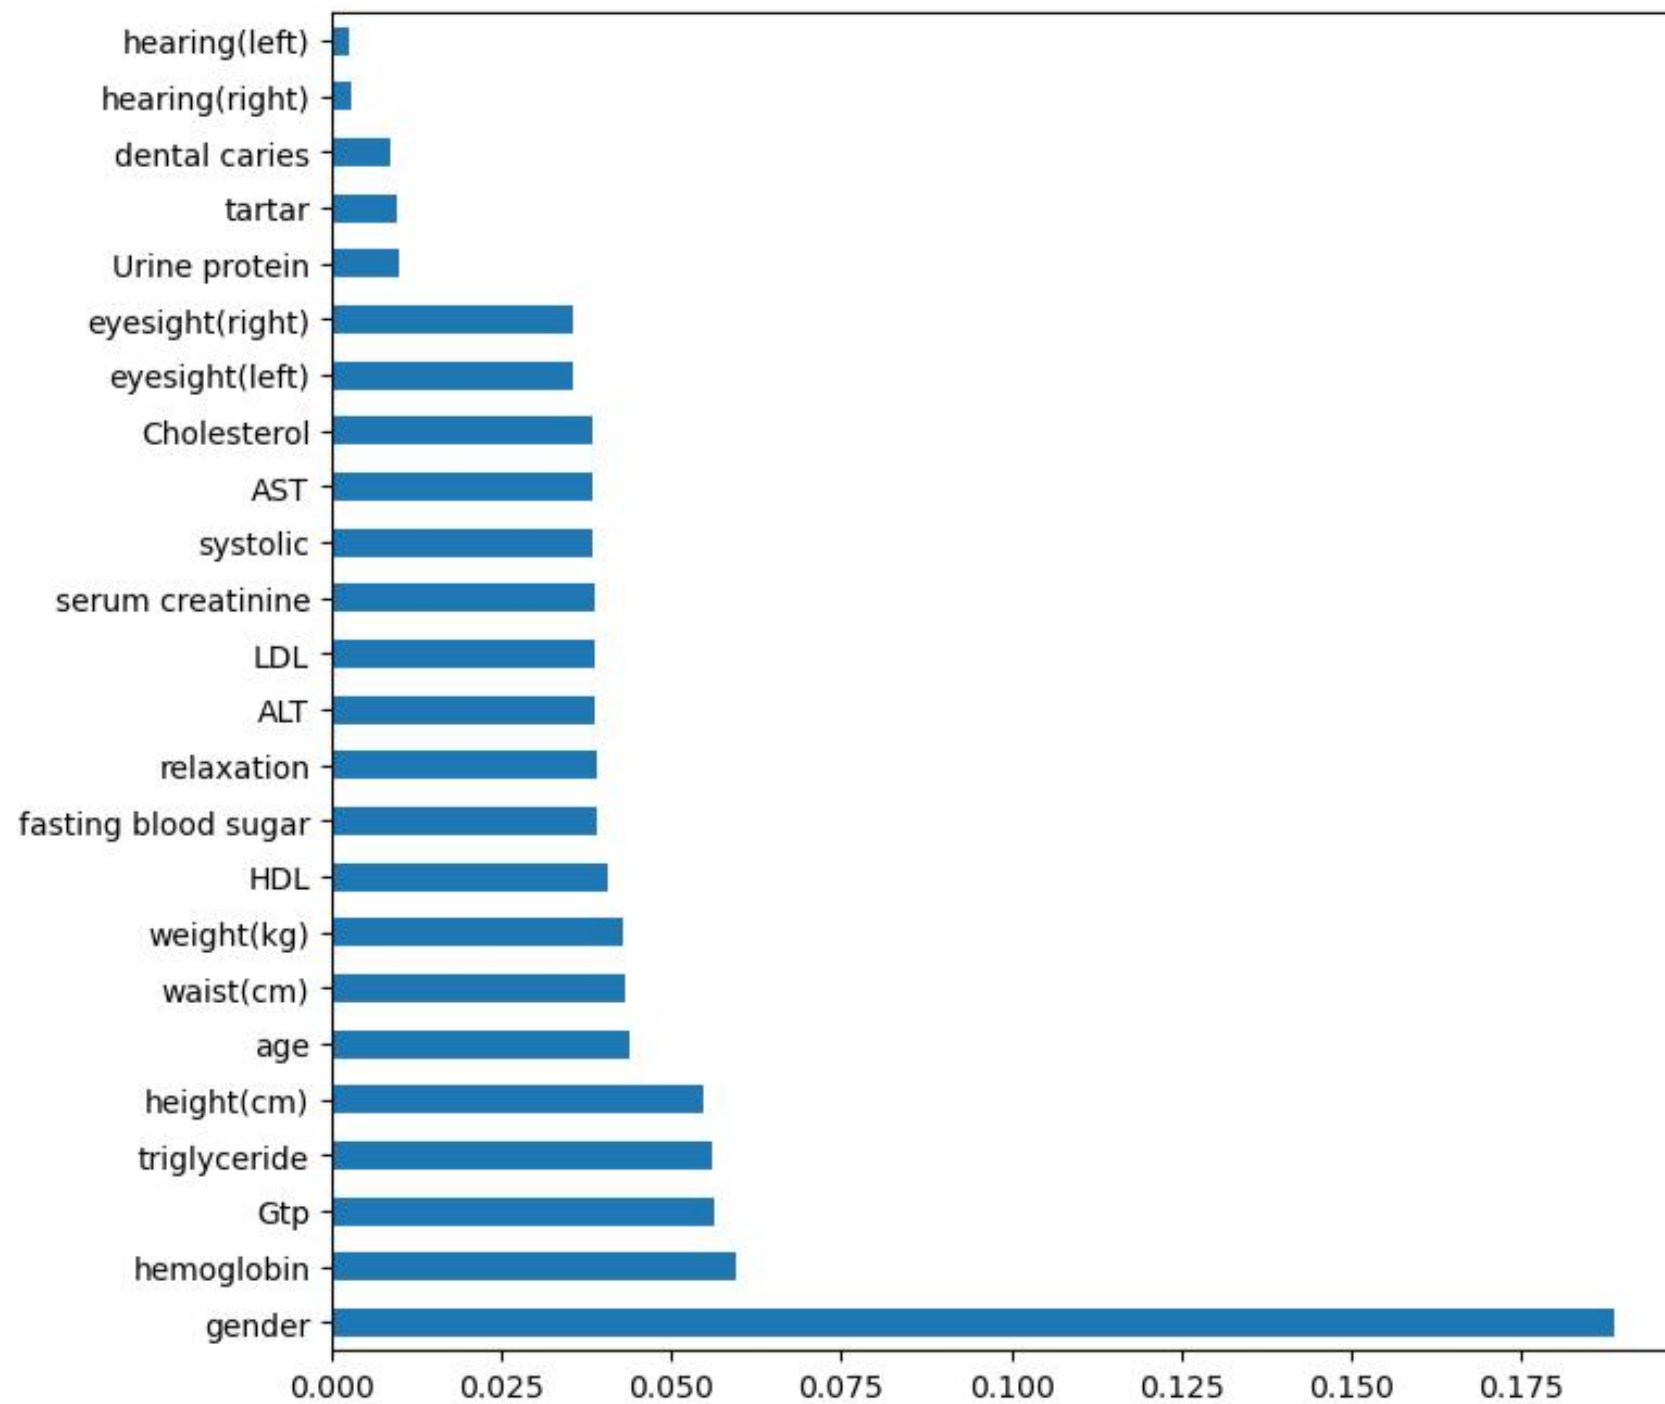
In [25]: *#Feature importance is a technique that calculate a score for all the input features for a given model.
#So out of 24 features we will select the top 15 features based on the score.*

```
X=df.iloc[:, :-1]
y=df["smoking"]

from sklearn.ensemble import ExtraTreesClassifier
model=ExtraTreesClassifier()

model. fit(X,y)
df1=pd.Series(model.feature_importances_,index= X.columns)
plt.figure(figsize=(8,8))
df1.nlargest(24).plot(kind="barh")

plt.show()
```



Logistic Regression



```
n [30]: #Calculating accuracy and generating the classification report of Logistic Regression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report

# Define X and y
X = df[['gender', 'height(cm)', 'Gtp', 'hemoglobin', 'triglyceride', 'age', 'weight(kg)', 'waist(cm)', 'HDL', 'serum creatinine',
y = df['smoking']]

# Split the data into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale the data using StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)

# Train a Logistic regression model
lr = LogisticRegression()
lr.fit(x_train, y_train)

# Make predictions on the test set
y_pred = lr.predict(x_test)

# Evaluate the model's accuracy and generate a classification report
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')
print(report)
```

Accuracy: 0.73

	precision	recall	f1-score	support
0	0.81	0.76	0.78	7027
1	0.63	0.69	0.66	4112
accuracy			0.73	11139
macro avg	0.72	0.73	0.72	11139
weighted avg	0.74	0.73	0.74	11139



Decision Tree



In [33]: *#The accuracy of the logistic regression model is 78 percentage*

```
from sklearn.tree import DecisionTreeClassifier
# Train a decision tree classifier
dt = DecisionTreeClassifier()
dt.fit(x_train, y_train)

# Make predictions on the test set
y_pred = dt.predict(x_test)

# Generate a classification report
report = classification_report(y_test, y_pred)

print(report)
```

	precision	recall	f1-score	support
0	0.83	0.82	0.83	7027
1	0.70	0.71	0.70	4112
accuracy			0.78	11139
macro avg	0.76	0.76	0.76	11139
weighted avg	0.78	0.78	0.78	11139

Bagging Algorithm – Bagging Classifier

```
n [35]: #Bootstrap Aggregation or bagging involves taking multiple samples from the training dataset  
 #(with replacement) and training a model for each sample  
from sklearn.ensemble import BaggingClassifier  
  
# Train a bagging classifier with decision tree base estimator  
bagging_clf = BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=1000)  
bagging_clf.fit(x_train, y_train)  
  
# Make predictions on the test set  
y_pred = bagging_clf.predict(x_test)  
  
# Generate a classification report  
report = classification_report(y_test, y_pred)  
  
print(report)
```

	precision	recall	f1-score	support
0	0.88	0.85	0.86	7027
1	0.75	0.80	0.77	4112
accuracy			0.83	11139
macro avg	0.82	0.82	0.82	11139
weighted avg	0.83	0.83	0.83	11139

Bagging Algorithm – Extra Trees ✓

```
In [36]: from sklearn.ensemble import ExtraTreesClassifier
# Train an Extra Trees classifier
et = ExtraTreesClassifier(n_estimators=1000, random_state=42)
et.fit(x_train, y_train)

# Make predictions on the test set
y_pred = et.predict(x_test)

# Generate a classification report
report = classification_report(y_test, y_pred)

print(report)
```

	precision	recall	f1-score	support
0	0.89	0.84	0.86	7027
1	0.75	0.82	0.78	4112
accuracy			0.83	11139
macro avg	0.82	0.83	0.82	11139
weighted avg	0.83	0.83	0.83	11139

Bagging Algorithm – Random Forest

In [37]: `from sklearn.ensemble import RandomForestClassifier`

```
# Train a random forest classifier
rfc = RandomForestClassifier(n_estimators=1060)
rfc.fit(x_train, y_train)

# Make predictions on the test set
y_pred = rfc.predict(x_test)

# Generate a classification report
report = classification_report(y_test, y_pred)

print(report)
```

	precision	recall	f1-score	support
0	0.88	0.84	0.86	7027
1	0.75	0.80	0.78	4112
accuracy			0.83	11139
macro avg	0.82	0.82	0.82	11139
weighted avg	0.83	0.83	0.83	11139
